

Hadoop

Yizheng (Ethan) Chen

Advisor: Prof. Aditya Akella

Outline



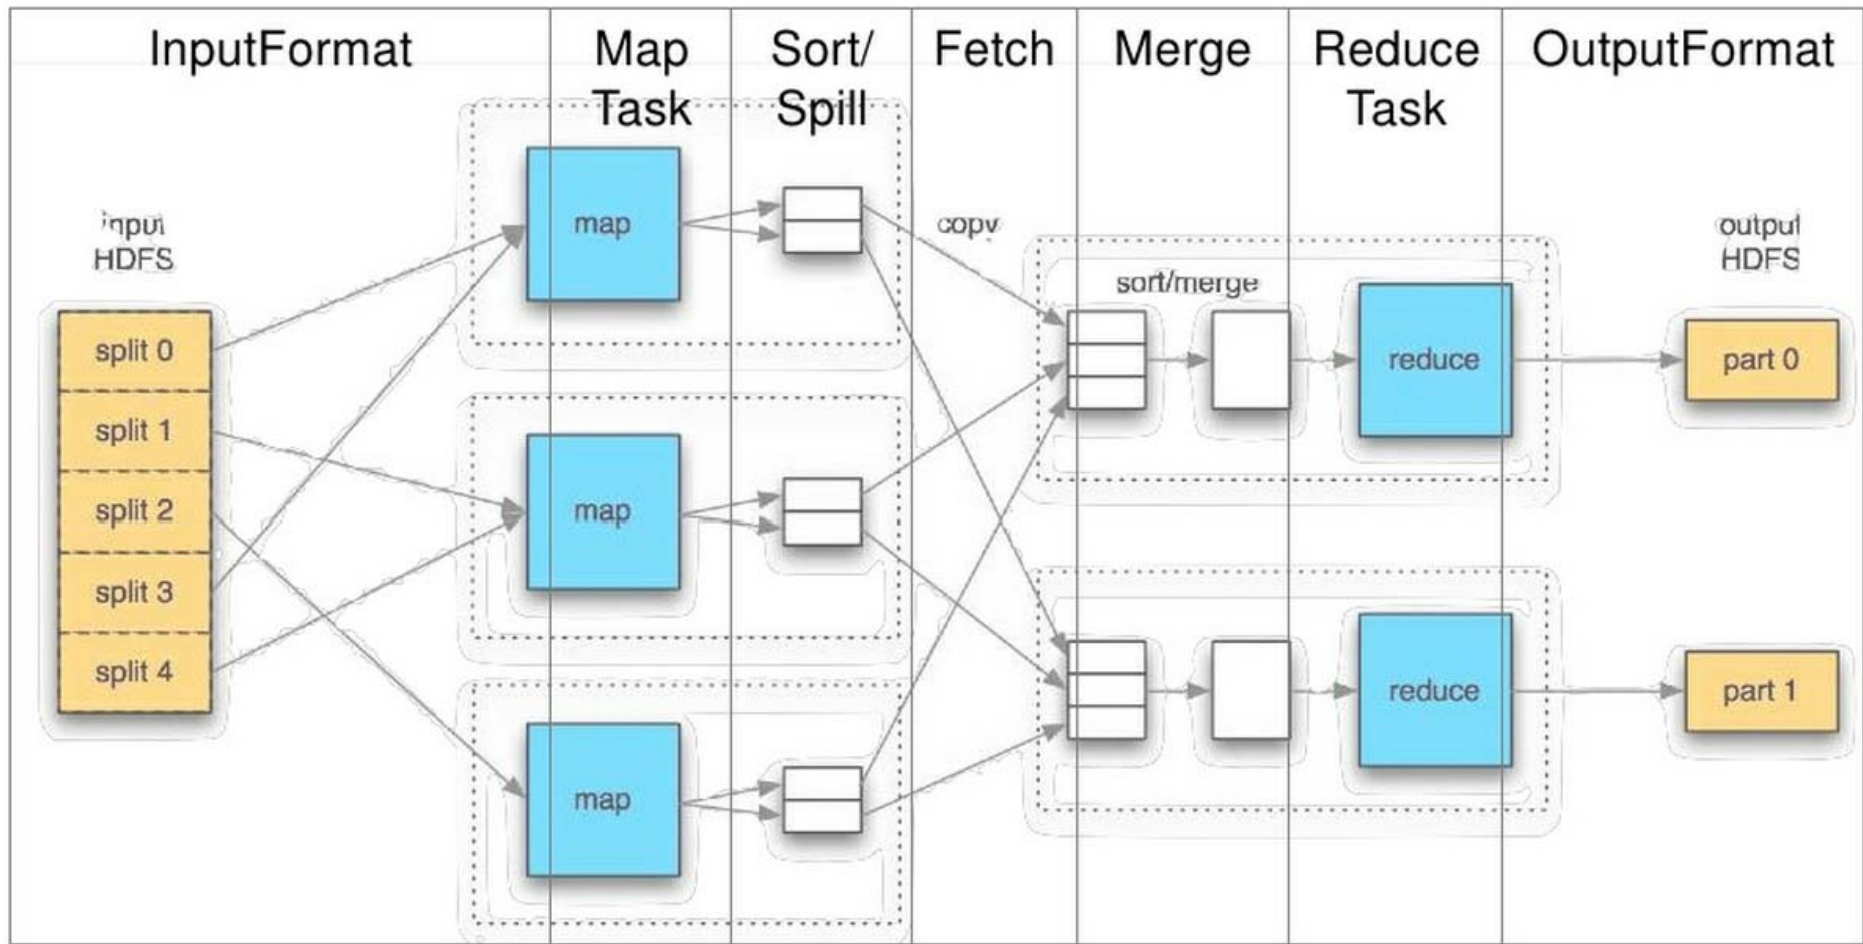
Hadoop

Yarn (NextGen Hadoop)

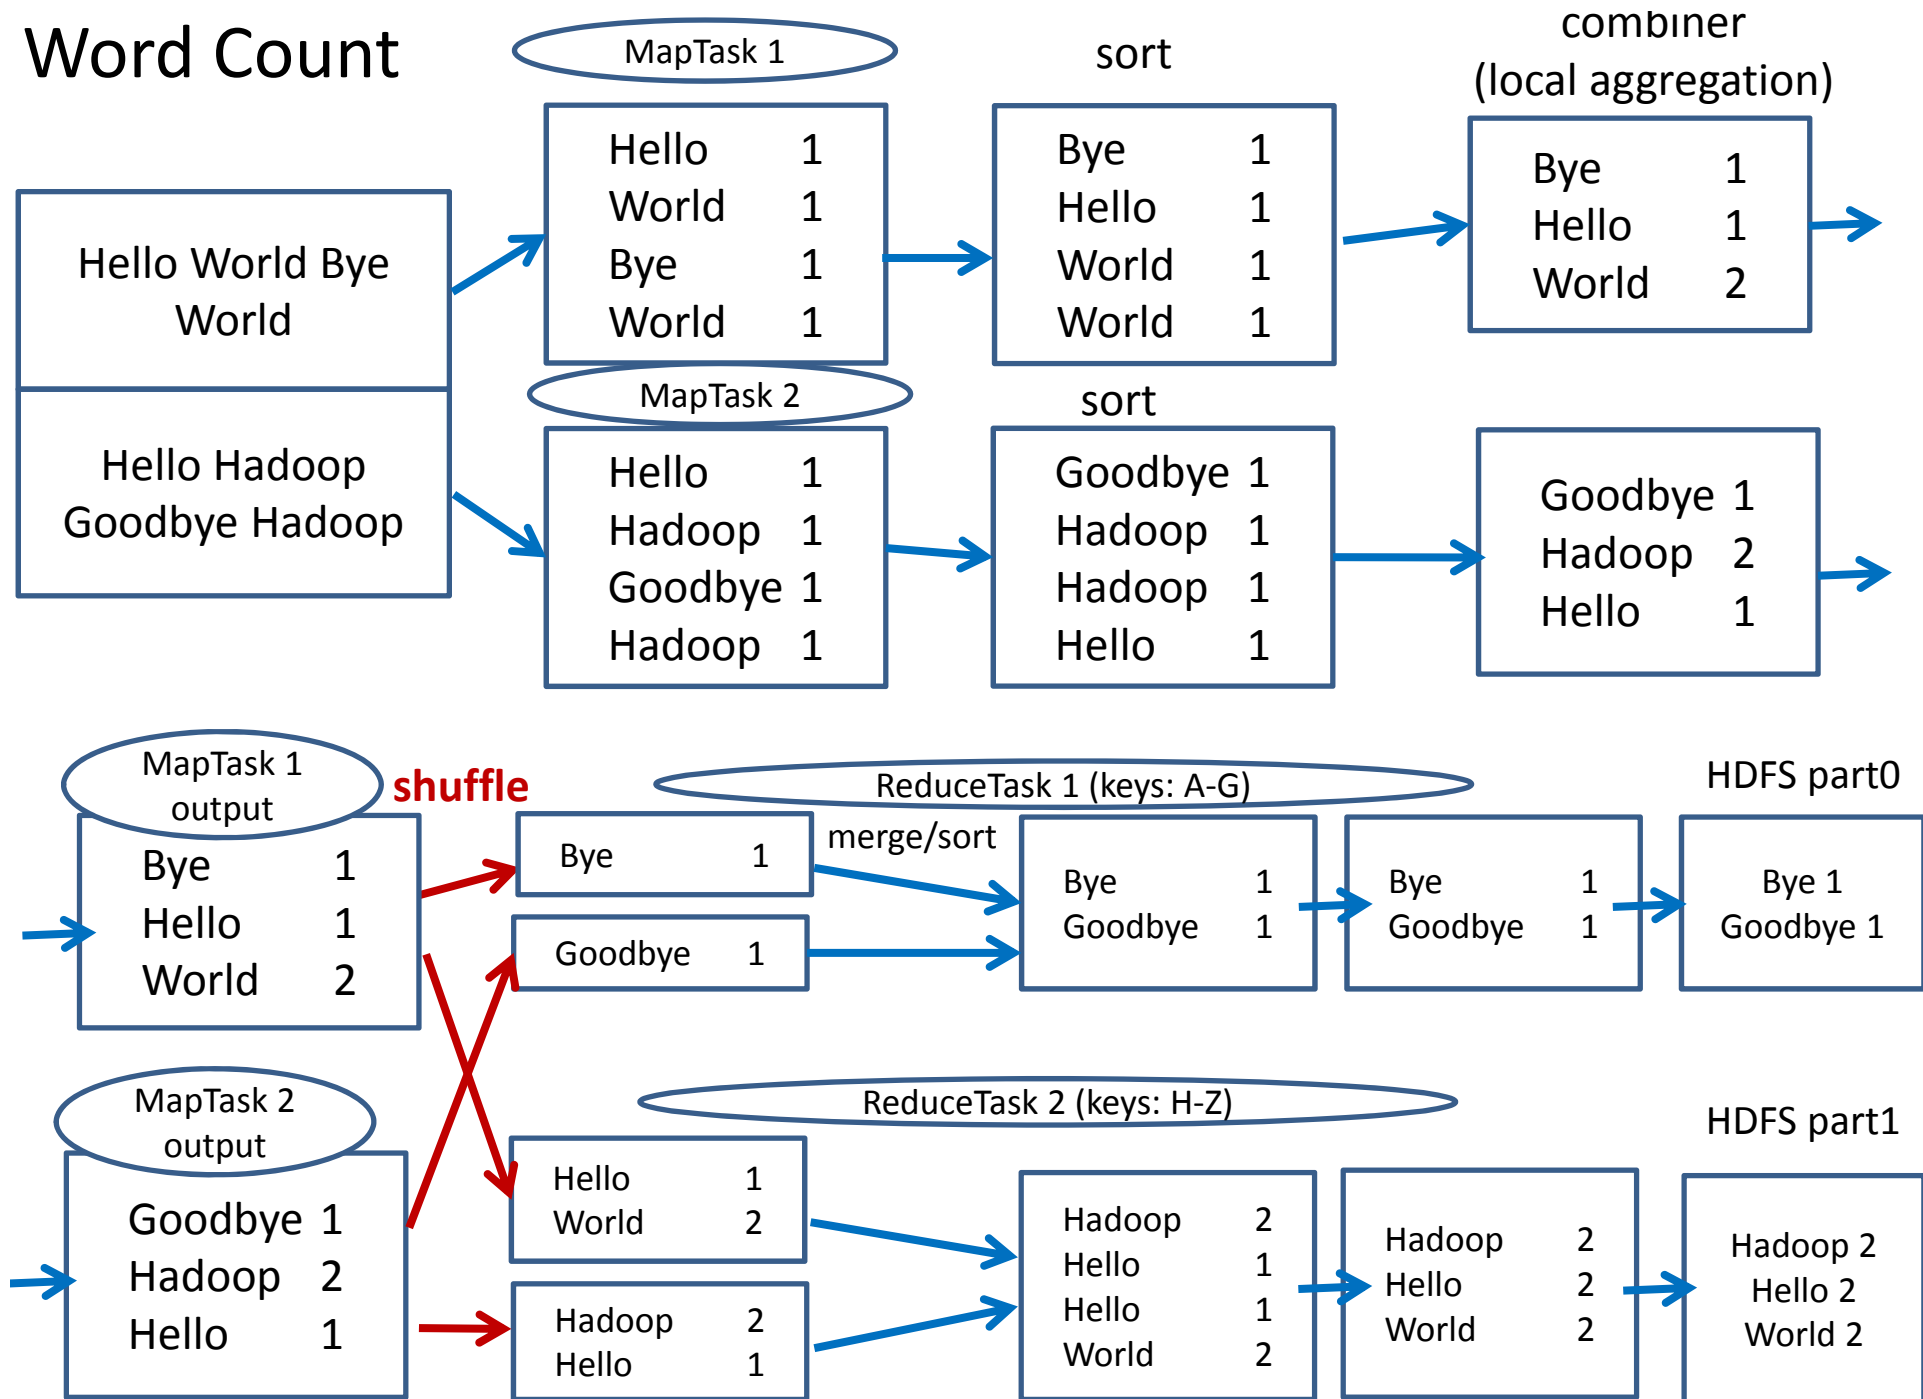
Hadoop

- What is Apache Hadoop
 - A framework (open-source software) for reliable, scalable, distributed computing
 - Hadoop MapReduce
 - A system for parallel processing of large data sets
 - Hadoop Distributed File System (HDFS™)
 - A distributed file system that provides high-throughput access to application data
 - Similar to GFS
 - <http://hadoop.apache.org/>
- Why Hadoop?

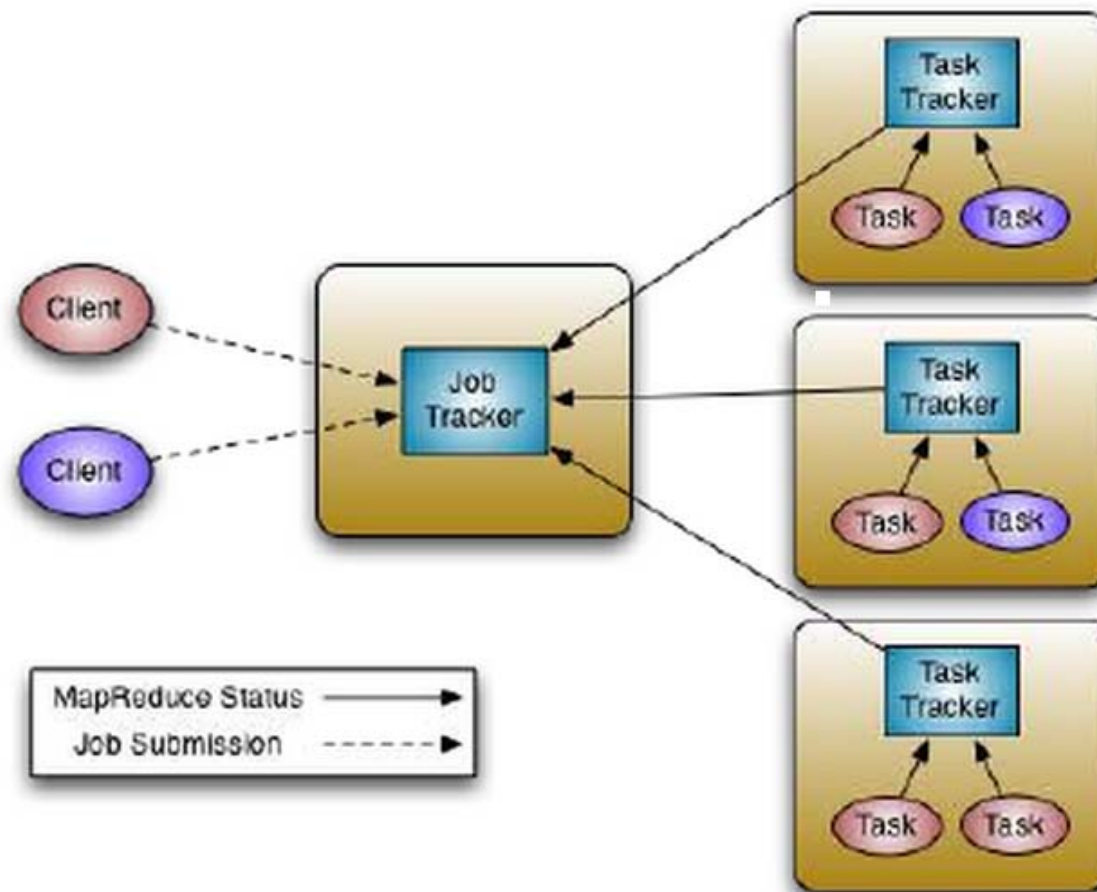
Hadoop Job Execution



Word Count



Hadoop MapReduce




Hadoop Schedulers

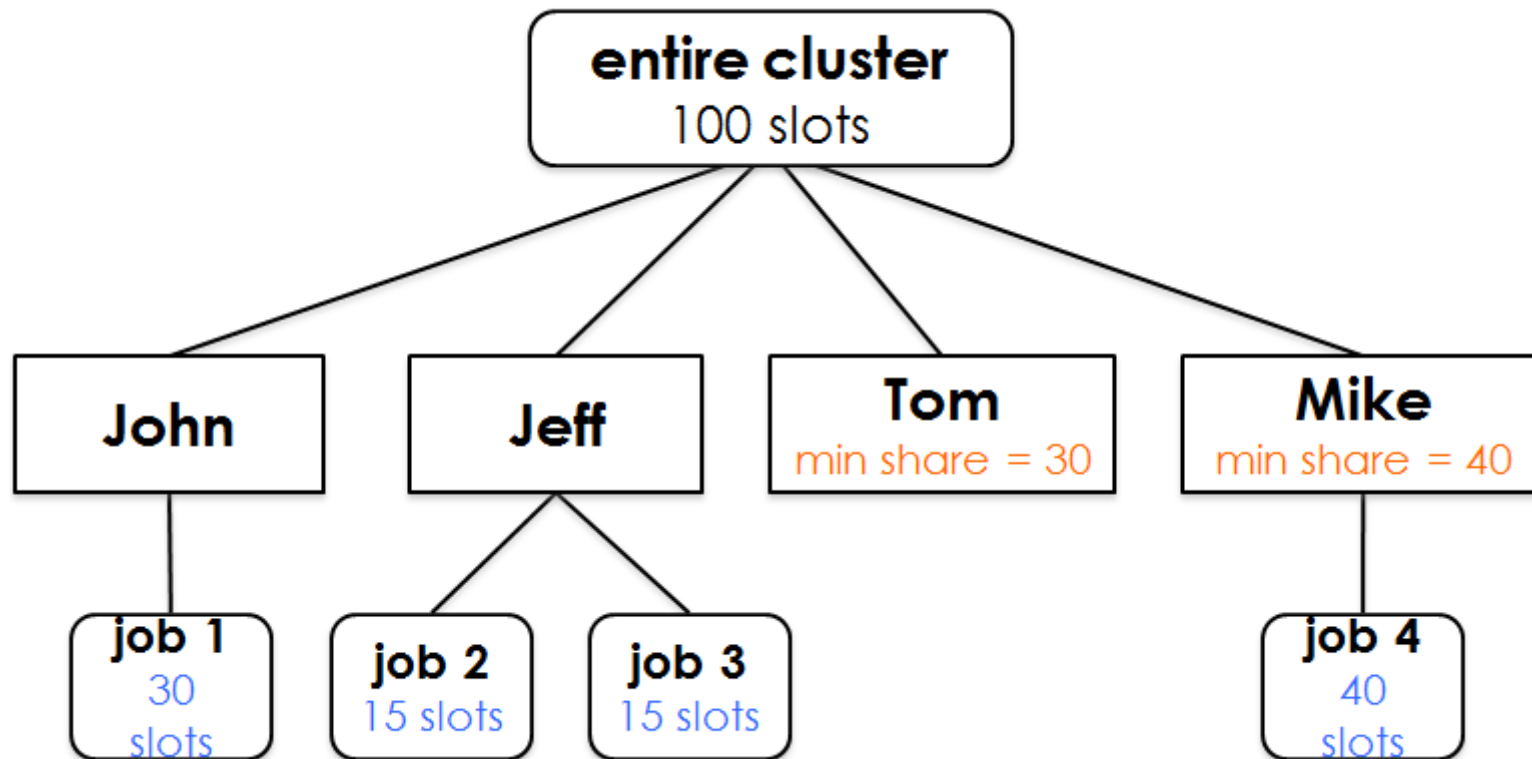
- A pluggable framework for job scheduling algorithm available since Hadoop 0.19
 - FIFO
 - Fair Scheduler (Facebook)
 - Capacity Scheduler (Yahoo!)

FIFO scheduler

Originally optimized for large batch jobs(web index construction)
FIFO order + priority queues

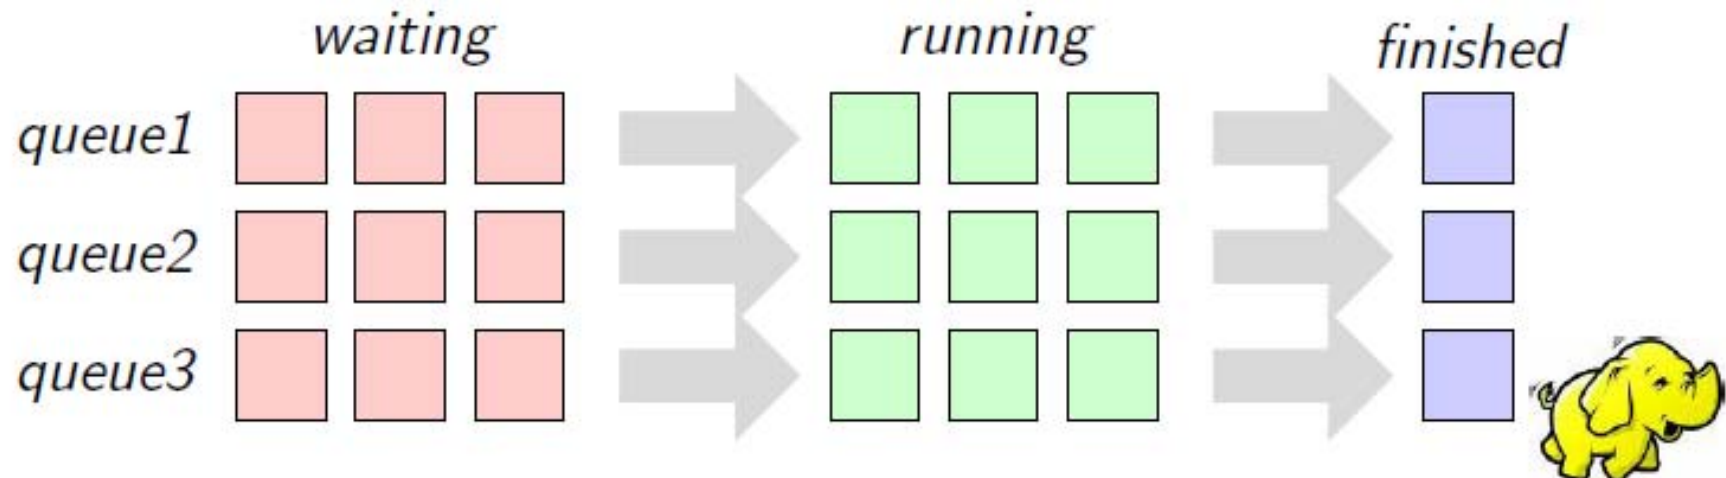


Fair Scheduler

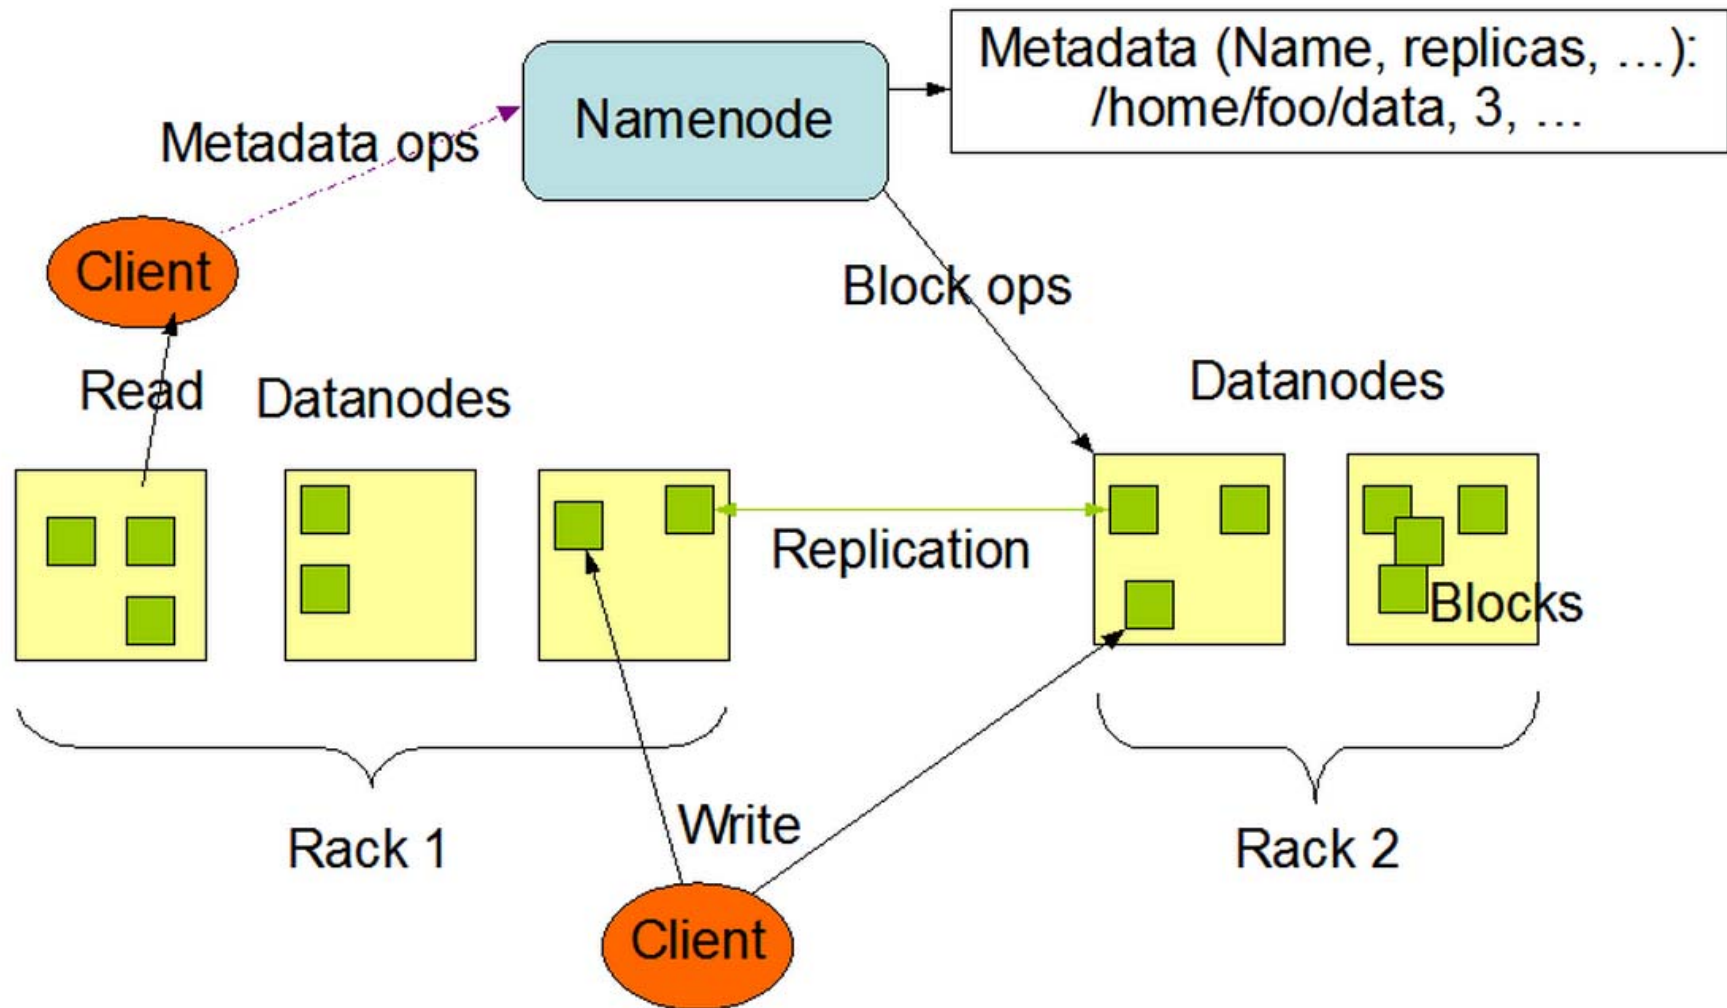


Capacity Scheduler

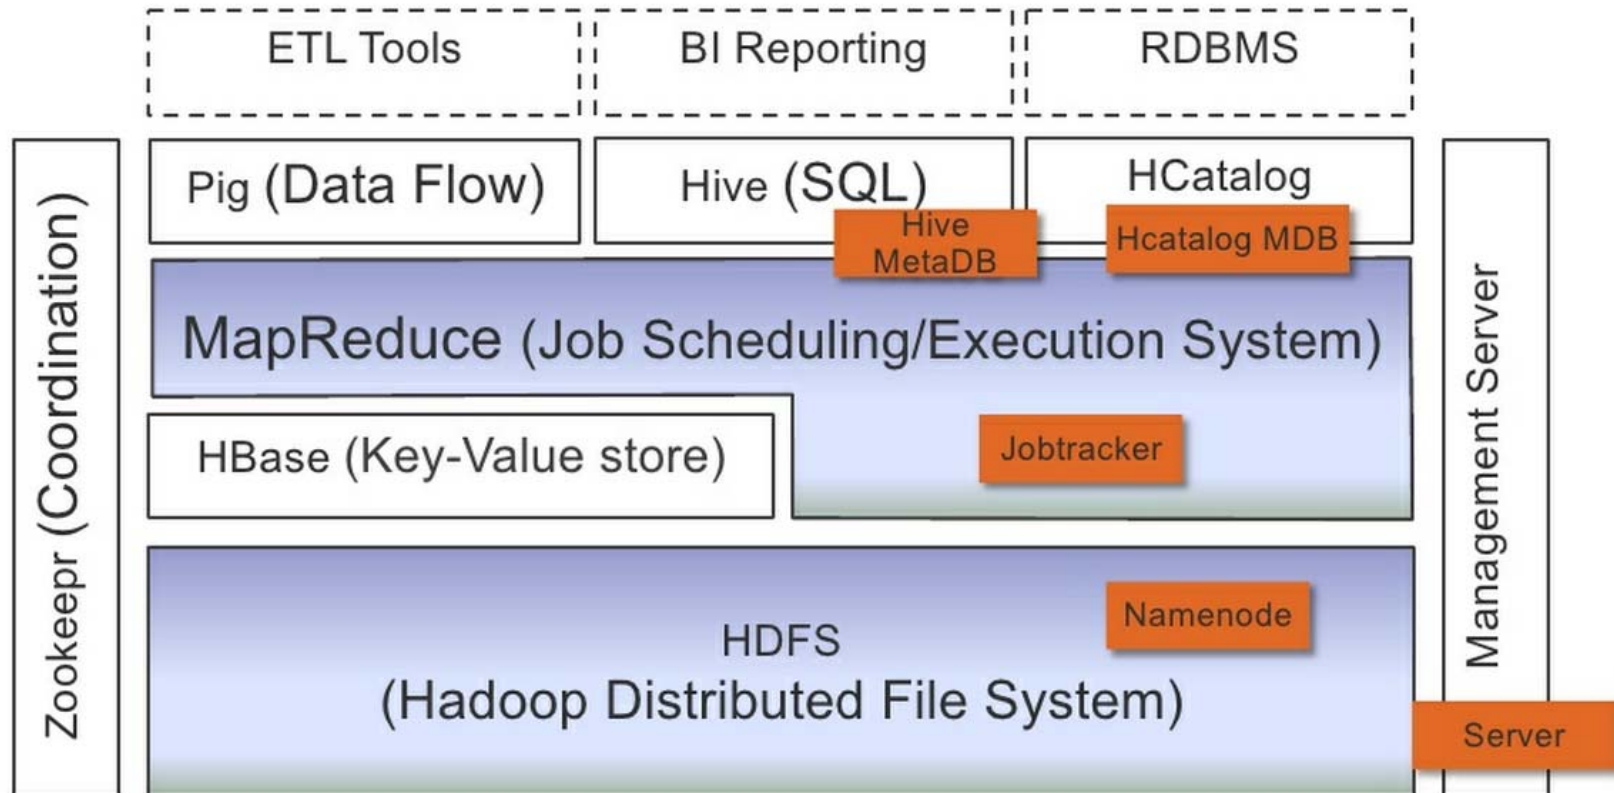
- Organizes jobs into queues
- Queue shares as %'s of cluster
 - Optionally, can limit maximum resources per queue
- FIFO scheduling within each queue
- Contributed by Yahoo!



HDFS Architecture



Hadoop Ecosystem



HBase: BigTable-like

Hive: Data summarization and ad hoc querying

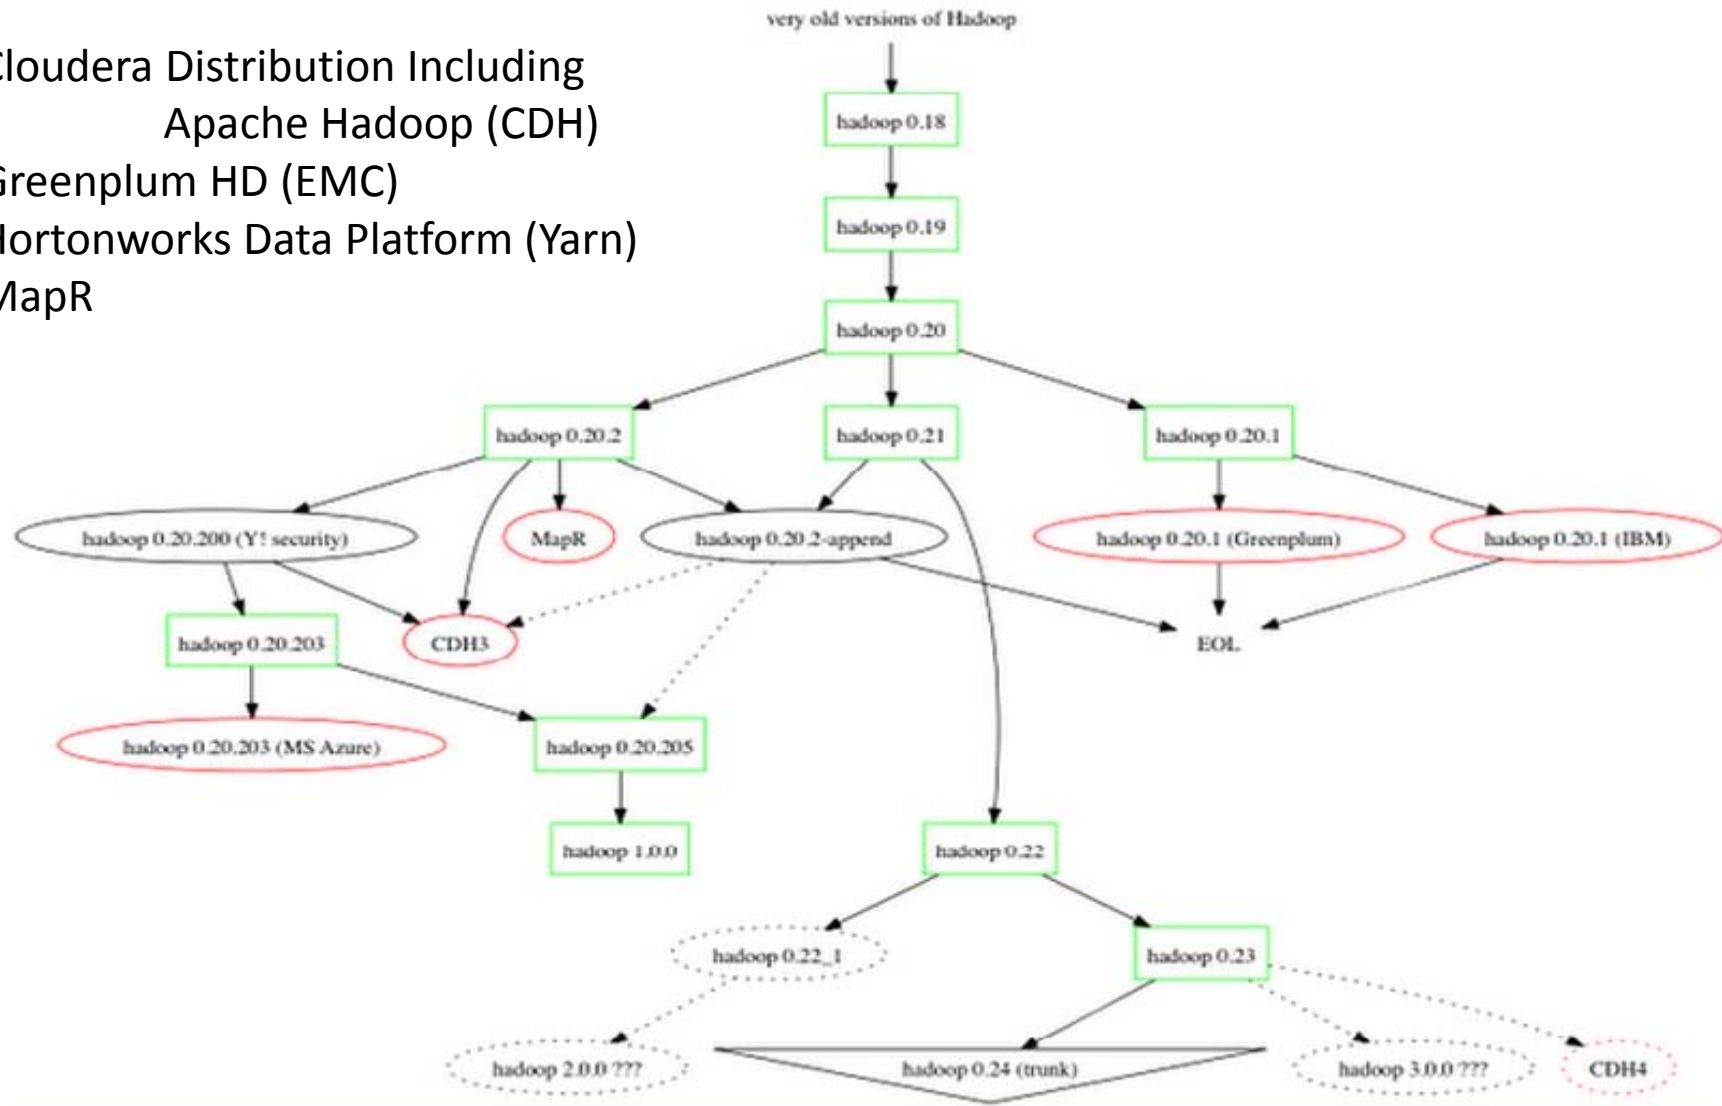
Pig: A high-level data-flow language and execution framework for parallel computation

HCatalog: Table and storage management service (table abstraction of data)

Zookeeper: A high-performance coordination service for distributed applications

Hadoop and Hadoop-derived Distributions

- *Cloudera Distribution Including Apache Hadoop (CDH)
- *Greenplum HD (EMC)
- *Hortonworks Data Platform (Yarn)
- *MapR



https://blogs.apache.org/bigtop/entry/all_you_wanted_to_know

Outline



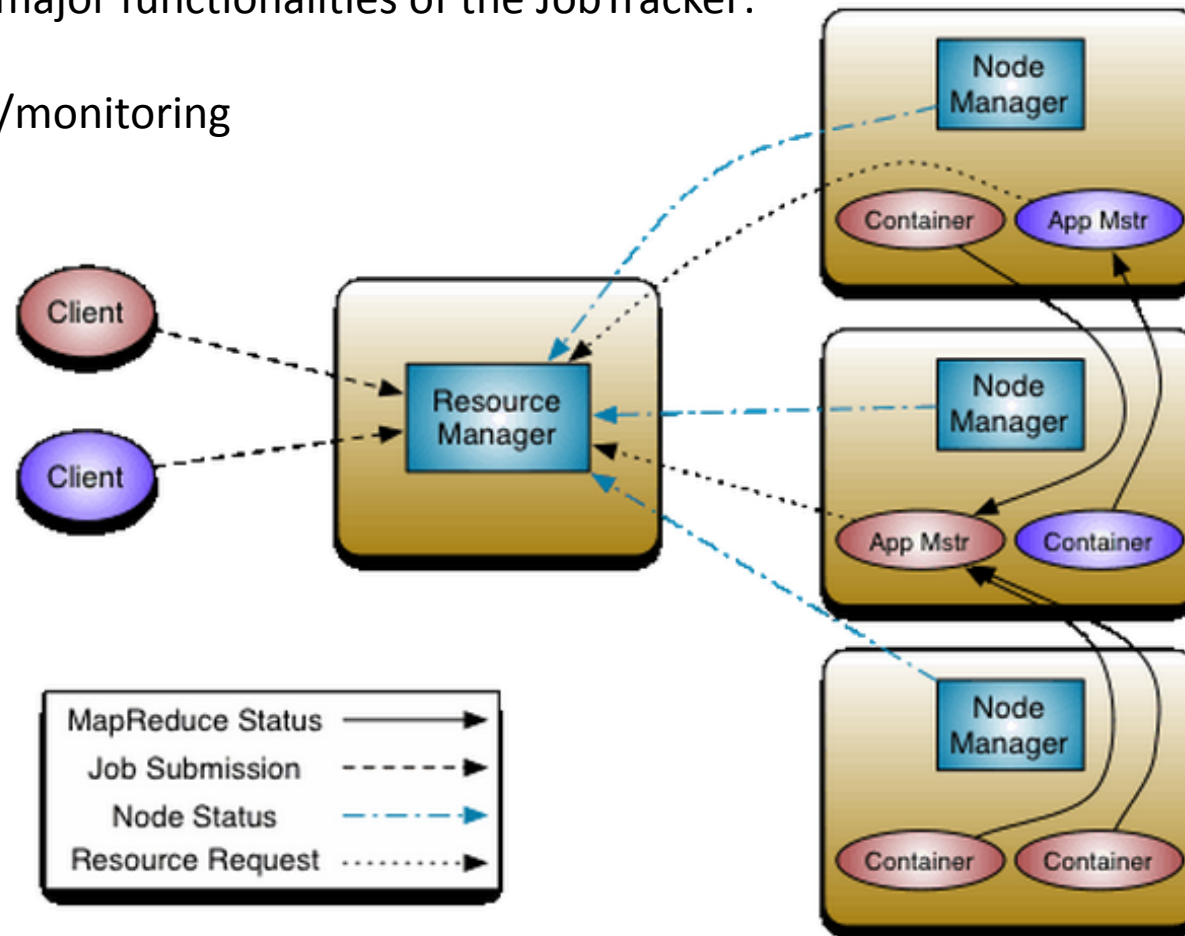
Hadoop

Yarn (NextGen Hadoop)

Yarn (NextGen Hadoop)

Split up the two major functionalities of the JobTracker:

- * Management
- * Job scheduling/monitoring



ResourceManager:

- *Scheduler: allocate resources to the various running applications (pluggable policy plug-in)
- *ApplicationsManager : accept job-submissions/launch the first container for ApplicationMaster

Resource Allocation

- the resource request understood by the Scheduler is of the form:

<priority, (hostname/rackname/*), capability, #containers>

- **Scheduler API**

There is a single API between the Scheduler and the ApplicationMaster:

(List <Container> newContainers, List <ContainerStatus> containerStatuses) allocate (List <ResourceRequest> ask, List<Container> release)

The input-splits are:

task-id	split-0	split-1	split-2
map0	r11/h1001	r22/h2121	r31/h3118
map1	r11/h1010	r22/h2121	r45/h4123
reduce0	*		

The corresponding resource requests from the ApplicationMaster to the Scheduler:

Priority / Capability	Hostname	#Containers
2 / 1GB	h1001	1
	h1010	1
	h2121	2
	h3118	1
	h4123	1
	r11	2
	r22	2
	r31	1
	r45	1
	*	2
1 / 1GB	*	1

Compact:
 $O(\text{clustersize})$

If a container is available on `r22/h2121`, the Scheduler allocates it to the application and invalidates requests on `h2121`, `r22` and `*` by updating the resource request list as follows:

task-id	split-0	split-1	split-2
map0	r11/h1001	r22/h2121	r31/h3118
map1	r11/h1010	r22/h2121	r45/h4123
reduce0	*		

Priority/ Capability	Hostname	#Containers
2 / 1GB	h1001	1
	h1010	1
	h2121	2 1
	h3118	1
	h4123	1
	r11	2
	r22	2 1
	r31	1
	r45	1
	*	2 1
1 / 2GB	*	1

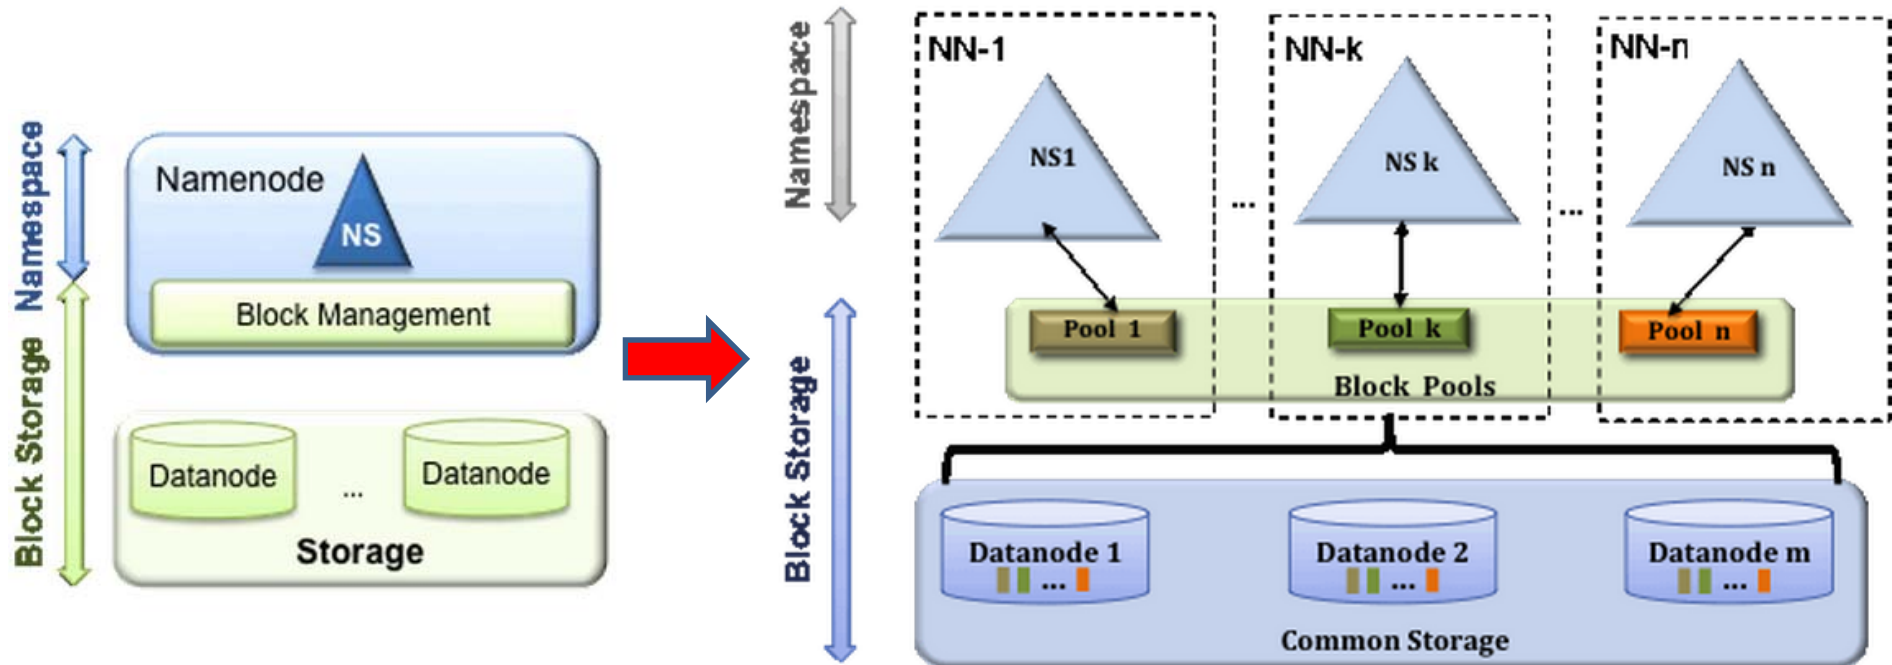
The ApplicationMaster can then, optionally, invalidate requests for h1001, h3118, r11 and r31 after it decides to run *map0* on h2121.

Now, if a container is available on r11/h1010, the Scheduler allocates it to the application and invalidates requests on h1010, r11 and * as follows:

task-id	split-0	split-1	split-2
map0	r11/h1001	r22/h2121	r31/h3118
map1	r11/h1010	r22/h2121	r45/h4123
reduce0	*		

Priority / Capability	Hostname	#Containers
2 / 1GB	h1001	1
	h1010	1 0
	h2121	1
	h3118	1
	h4123	1
	r11	1 0
	r22	1
	r31	1
	r45	1
	*	1 0
1 / 1GB	*	1

HDFS Federation



Benefits:

- * Namespace Scalability
- * Performance
- * Isolation

Q & A